

Trust and Transparency in Generative AI: Comparing Perceptions of Commercial and Institutional Platforms in a University Population

Sasha Watanabe
University of Florida

Sara Waters
University of Florida

Abstract

Generative Artificial Intelligence (AI) and the rise of institutionally managed systems are transforming higher education. However, little is known about how the university community perceives the privacy and security of university generative AI tools, particularly compared to commercial platforms. We base our decision to evaluate perceptions of privacy and security on the inverse relationship between trust and perceived risk. We surveyed 57 students, faculty, and staff at the University of Florida (UF) and conducted three supporting interviews to investigate (1) how these groups use generative AI tools, (2) how they perceive privacy and security in commercial versus institutional platforms, namely NaviGator Chat, and (3) what risks arise from limited transparency regarding data handling and policies. Our results indicate that the university community uses generative AI extensively, but behavior and trust are motivated by peers and assumptions. Students use NaviGator Chat less often than commercial platforms; yet, UF’s institutional status elicits a slightly higher level of trust. We also show that although users attempt to manage their privacy, their depth of understanding often constrains effective self-protection. These results guide our recommendations for deploying transparent and ethical generative AI in higher education.

1 Introduction

Generative Artificial Intelligence (AI) has become increasingly embedded in higher education and workplaces, where it is reshaping teaching, research, and learning. Large-language models (LLMs) such as OpenAI ChatGPT, Google Gemini, and Anthropic Claude provide a multitude of benefits in various contexts, from business and technology to healthcare and society [17] [19]. As these platforms grow in popularity, questions about their behavior, policies, and perceptions—many of which remain unanswered and underexplored—also increase [22] [18]. In July 2024, the University of Florida (UF) launched a new institutional system dubbed NaviGa-

tor Chat [6]. NaviGator Chat promises a reliable, private, and safe way to utilize the best LLM chatbot models for academic and professional work [12]. However, no formal evaluation has documented how much the university community trusts NaviGator Chat, nor how they perceive its privacy and security compared to commercial alternatives. Therefore, the security and privacy challenges posed to commercial platforms are equally relevant to NaviGator Chat.

Trust has long linked users’ perceived reliability of a system with their willingness to disclose information [11]. Recent studies distinguish between *human-like* factors, such as benevolence and integrity, and *system-like* factors, such as reliability and functionality. Some scholars argue that trust is induced by the perceived *humanness* of technology, whereas others proclaim that trust reduces privacy concerns and increases self-disclosure [8] [9]. Consequently, popularity-driven trust is particularly concerning if a system is popular but unethical [20]. Institutionally hosted platforms such as NaviGator Chat can demonstrate transparent, ethical data-handling practices and educate their communities to better understand privacy and security when using generative AI.

This opportunity is especially critical because the UF student body, by virtue of their age and digital literacy, is in a vulnerable position. Students regularly externalize personal and academic data online without fully understanding how their information might be stored, sold, or reused [7]. Since UF does not mandate courses on technology and AI literacy—unlike courses on violence prevention and mental health—the university presumes that students possess some standard of knowledge on these topics. This raises the question: how well do UF students understand what data is sensitive, how to protect their digital footprint, and how to interact safely with technology and generative AI systems?

Students who use non-transparent generative AI tools may inadvertently expose sensitive information without proper awareness of privacy risks. While NaviGator Chat labels select models as secure for sensitive and restricted data, the system’s protocol for logging and managing chats remains opaque. Furthermore, although the NaviGator Chat data labels

are based on UF’s data labels, the NaviGator Chat interface does not explain their function. [13]. When a new chat is started, users can click on the list of data labels associated with the model to read about their meanings, but the text is small and easily overlooked. Without the ability to understand how NaviGator Chat upholds its promises of safety and privacy, the platform magnifies the risk of inadvertent misuse or low engagement despite its potential benefits.

Therefore, the purpose of this study is threefold: (1) quantify how extensively UF students, faculty, and staff use generative AI tools and which ones; (2) investigate how these populations perceive the privacy and security of generative AI tools, with a focus on comparing commercial platforms with institutional platforms like NaviGator Chat; and (3) identify risks that arise through how participants understand data-handling mechanisms and policies to formulate design principles for responsible and ethical generative AI systems in higher education.

The methodology employs a mixed-methods design consisting of a literature review, survey, and semi-structured interviews with UF students, faculty, and staff. The online survey was distributed on Qualtrics to a broad sample of UF students, faculty, and staff to collect quantitative data on platform usage, perceived privacy and security risks, competence, and behavioral adaptations to risks. Semi-structured interviews complement the survey by uncovering narratives, misconceptions, and technical details that shape these quantitative indices. Quantitative data was analyzed using descriptive and correlation statistics, while qualitative data was analyzed using thematic coding. Combining these methods provides a deeper understanding of how trust and privacy and security concerns translate into usage patterns and inform institutional policies.

Our results indicate that although only 10% of survey respondents stated that NaviGator Chat has somewhat worse privacy protection compared to commercial generative AI platforms, 27% reported disclosing more personal or sensitive information on commercial platforms, while 43% reported that their level of data disclosure is consistent across platforms. In fact, the average level of trust in NaviGator Chat’s data handling practices was marginally higher at 2.51 compared to ChatGPT at 2.33 and Microsoft Copilot at 2.22 (on a scale of 1 to 5, with 5 being complete trust). Survey respondents primarily evaluate generative AI platforms based on public reputation, accurate and unbiased output, and recommendations from family, friends, or peers. Moreover, 45% of survey respondents alter their behavior or avoid certain generative AI systems due to privacy or trust concerns. Only 28% of survey respondents reported ever using NaviGator Chat.

These findings suggest that despite the potential privacy and security benefits of NaviGator Chat, the UF community is generally unaware of its offerings and prefers commercial platforms, namely ChatGPT, due to their positive public perception and the trust of their peers. Nevertheless, these results

serve to inform the University of Florida and other academic institutions with similar platforms about how to develop their message, user experience, and policies to encourage the use of their generative AI tools and educate their communities about how to use generative AI safely, ethically, and transparently. Ultimately, this paper provides insights and actionable recommendations for building generative AI services that foster trust by protecting users’ privacy and data security.

2 Problem Specification

As generative AI integrates further into academic practice, questions of trust, privacy, and behavior become increasingly relevant. NaviGator Chat is described as “the first comprehensive, self-service AI platform in U.S. higher education,” according to the University of Florida [6]. However, NaviGator Chat’s usage and how its privacy implementation compares to the data-handling practices of commercial platforms are yet to be researched. This presents a unique opportunity to study experiences and approaches to using institutional generative AI tools.

According to UF’s Artificial Intelligence website, under Teaching with AI, NaviGator Chat provides access to leading LLMs through a private, UF-hosted interface designed for academic and professional work. User input remains on UF or UF-approved vendor servers until users delete it [12]. Although this page states that “NaviGator Chat is approved for open data only,” the interface lists models that feature three data labels: restricted, sensitive, and open data. Since UF seeks to foster “AI Across the Curriculum,” understanding how students, faculty, and staff use and perceive NaviGator Chat is crucial for improvements [16]. The inspiration for this project stems from a lack of conversation and use of NaviGator Chat by students, contrasted with an increase in training and use by staff.

Without officially published data-handling practices and policies, we referred to the NaviGator Chat model “gpt-oss-20b,” which is labeled safe for restricted data, to inquire how NaviGator Chat might process data. The pipeline begins when anyone with a GatorLink account enters a prompt into NaviGator Chat (the input method is typing or voice, which is transcribed into text). The prompt lands on a UF-hosted server where the input is fed to the chosen LLM deployed within UF data center containers. The model generates a response that is then returned to the user’s browser. The request-response pair is allegedly logged in a secure database for auditing and fine-tuning of the isolated model. Logs are kept for a predetermined retention period that is not publicly defined. Access to the data is apparently limited to UF staff and system administrators who are bound by UF privacy and data-use regulations. Although this pipeline has not yet been confirmed, NaviGator Chat claims that user data is not shared with external parties, including commercial AI vendors.

2.1 Research Questions

The primary objective of this project is to investigate the attitudes of UF students, faculty, and staff toward generative AI tools, how these attitudes are formed, and how they impact the use of generative AI platforms. Specifically, this project seeks to ask the following research questions:

RQ 1: How extensively do UF students, faculty, and staff use generative AI, and which tools do they use?

RQ 2: To what extent do people trust generative AI systems to handle their data responsibly and securely, and is trust different between commercial and institutional systems? What factors most prominently shape these perceptions? How do these perceptions of trust influence decisions in the use of generative AI?

RQ 3: What do people understand about how generative AI platforms handle data and privacy? How do people perceive the privacy protections of commercial platforms versus the models supported by NaviGator Chat, especially those labeled for sensitive data?

RQ 4: What are the privacy risks for populations within the university community in relation to generative AI systems? What should be considered when creating a generative AI system for higher education?

Addressing these questions contributes to ethical AI design, particularly in academia, to help institutions ensure the transparency and trustworthiness of their custom generative AI tools. Although NaviGator Chat offers a contained, institutionally managed environment for popular LLMs, the current lack of transparency regarding its privacy policies and data management practices leaves students vulnerable to potential misuse through undisclosed data collection and storage of personal or academic data. Consequently, the project aims to identify patterns in perceptions of privacy and trust when interacting with a variety of generative AI tools by producing a comparative analysis between UF’s NaviGator Chat and commercial AI platforms and provide actionable recommendations for ethical and privacy-conscious AI systems in academia.

3 Related Work

We first establish a theoretical foundation by conducting a literature review on attitudes toward privacy, trust in technology, and ethical AI practices in higher education. Through this review, we examine previous research on AI ethics, data protection, and credibility, as well as identify gaps in existing scholarship. The literature review also evaluates UF policies related to NaviGator AI for clarity, ethics, and transparency—much of which is not yet publicly addressed.

Attitudes toward generative AI have been broadly studied in recent years. Mahmoud et al. (2025) analyzed sentiments toward generative AI from YouTube comments using big data methods, revealing a generally negative sentiment re-

garding generative AI and identifying a variety of sentiment themes, one of which was privacy decay [10]. Technical fear and job threats were found to be factors in generative AI attitudes among employees in the hospitality and tourism industry, which is consistent with Mahmoud et al.’s YouTube study [25]. We investigate whether these themes also appear among the UF population, whether they differ in prevalence by demographic data, and how they might relate to privacy concerns.

Previous studies on trust in AI systems among students often define trust based on output correctness, considering the risks of hallucination or misinformation, not privacy [1]. Wang et al. (2025) conducted semi-structured interviews with university students based on the Antecedents-Privacy Concerns-Outcomes (APCO) macro-model to understand privacy risks, how they are informed, and how privacy risks manifest in behaviors regarding generative AI [21]. The APCO model and interview questions were relevant for this project and helped guide our interview methodology, which we modified for our larger population of students, faculty, and staff.

Researchers have proposed policy and guidance frameworks for adopting generative AI in higher education. Dotan et al. found that one of the barriers to adopting generative AI is lack of experience using it in higher education contexts; over half of surveyed faculty reported not using generative AI for professional duties [4]. In comparison, Amoozadeh et al. found that a large majority of surveyed university students reported using generative AI for academic purposes [1]. This research is particularly relevant for our study of generative AI usage among students, faculty, and staff at UF when AI guidance and policy have, albeit recently, been established with NaviGator Chat’s release. Golda et al.’s map of generative AI security and privacy concerns from five perspectives is also a helpful lens through which we examine privacy risks and mitigations in higher education and recommend how to administer policy and guidance on generative AI [5]. Golda et al. specifically raise concerns about inadequate data anonymization policies and the resulting surveillance risks. If NaviGator Chat links chat history with users’ identity, students may believe that the university is surveilling generative AI chats for plagiarism.

Recent research has also examined how the technical and design affordances of generative AI platforms shape user risk perception. Haoyu Zhao et al. applied affordance theory to explore how perceived interactivity, transparency, and agency impact user concerns across six dimensions: information, security, technical, social, ethical, and legal [24]. These researchers found that transparency and controllability significantly reduced perceived risks. Notably, women reported a higher perceived risk, an insight that helped inform our approach to understanding design cues and institutional context. As such, our survey collected demographic data—if respondents were comfortable providing these details—to investigate whether there are certain patterns between gender and usage

and perception of generative AI.

Another study by Yaman Yu et al. expanded the demographic lens to youth and categorized AI-related incidents into four groups: escalating mutual harm, intrapersonal harm, interpersonal harm, and autonomous generative AI harm [23]. The framework illustrates how privacy and well-being risks emerged through a feedback loop between users and generative AI systems. Although the study did not discuss these findings in the context of higher education, understanding how different user populations navigate privacy and self-disclosure remains important.

In the realm of academia, the use of ChatGPT and other large language models has begun to raise ethical and behavioral questions. A 2024 survey found that generative AI saves time but negatively impacts academic integrity [2]. The study emphasized the need for institutional policies that balance AI innovation with ethical responsibilities and transparency rather than ban AI tools. These findings are highly relevant to NaviGator Chat because its purpose is to provide a safe environment for professional and academic use. We must understand how trust and integrity concerns affect how the UF community interacts with generative AI in order to evaluate and improve generative AI management.

Additionally, UF’s policies [14] and media articles from *The Alligator* [3] and *Alachua County Today* [6] highlight the awareness and need for ethical AI practices but lack transparency in how UF handles data and privacy procedures specific to NaviGator AI. The gap in UF policy and media coverage further motivated our investigation to offer recommendations for clarity.

4 Design and Implementation

4.1 Research Hypothesis

University of Florida students, faculty, and staff perceive NaviGator Chat as more privacy-conscious and trustworthy than commercial AI platforms because it is customized for and managed by the institution. However, they may limit their input and interactions due to uncertainty about data handling, institutional oversight, and transparency regarding their privacy practices. These perceptions—shaped by factors such as prior experience with commercial AI platforms, institutional credibility, and awareness of privacy policies—influence how users disclose information and interact with NaviGator Chat compared to commercial generative AI platforms.

4.2 Methodology

To understand how the university community perceives privacy and security in generative AI, we utilized a mixed-methods approach of quantitative and qualitative research techniques. We compared UF’s institutional platform, NaviGator Chat, to commercial platforms such as OpenAI

ChatGPT, Google Gemini, and Anthropic Claude. Although NaviGator Chat hosts numerous popular commercial LLMs in a secure container, our initial understanding is that the UF community is generally unaware of what NaviGator Chat is and the benefits it offers—especially to students. Therefore, our goal is to understand how attitudes toward commercial versus institutional platforms influence user behavior, identify related risks, and verify the population’s literacy and awareness of generative AI.

4.2.1 Survey Study

A general survey was created and distributed to UF students, faculty, and staff to capture perceptions and experiences with generative AI. Through the survey, we collect quantitative data on how these populations understand privacy, trust, and information disclosures for commercial versus institutional AI systems, which offers insight into how these populations interpret privacy policies and how institutional branding shapes trust in generative AI systems. We administered the survey online through Qualtrics.

The survey included questions about demographics, generative AI usage habits, and privacy, with formats ranging from yes-no questions to multiple choice, Likert scales, and free response questions. All questions were optional and anonymous. Display logic was implemented on certain questions to prevent respondents from answering questions that were not applicable to them. A copy of all survey questions is printed in [Appendix A: Survey Questions](#).

The first section covered demographic questions such as their role at UF, degree program and major, department, gender, age, and race. Demographic data was used to help identify certain patterns specific to role, gender, or degree programs.

The second section covered generative AI use to discover how many respondents used generative AI chatbots, their familiarity with generative AI, and which generative AI platforms they frequently use. We also inquired about their typical use cases for generative AI. We asked participants whether they use NaviGator Chat, where they heard about it, and their use cases. Next, we asked NaviGator Chat users about their awareness and data input response to NaviGator Chat’s data labels (Open, Restricted, or Sensitive Data). Lastly, we asked what influenced their choice of generative AI platforms and whether they had encountered any difficulties while using such tools.

The final section was dedicated to privacy and security questions. We asked participants how much they trusted the following generative AI platforms: OpenAI ChatGPT, Google Gemini, Anthropic Claude, Microsoft Copilot, and UF NaviGator Chat. We also inquired about the transparency of each platform’s data handling practices, what privacy issues concern respondents, and whether they alter their behavior on generative AI platforms due to these concerns. Further questions were asked about mitigation practices, personal or sensitive

information disclosures to commercial platforms versus NaviGator Chat, and the extent to which respondents read policies or Terms of Service. The last question was an open-ended feedback and comments box.

Due to the limited time of the course during which this project was conducted, the survey was open for approximately three and a half weeks. We did not conduct a pilot, nor modify the survey as data populated. Since the survey primarily utilized binary, multiple-choice, and Likert-scale questions, the analysis was primarily quantitative. We received minimal responses to open-ended questions, such as “other” replies and the final feedback box. These answers are detailed in the [Participants](#) subsection (4.2.3) and the [Results](#) section (5).

4.2.2 Interview Process

Three semi-structured interviews were conducted in parallel with the survey to collect additional qualitative data points. The purpose of these interviews is to understand the reasoning behind attitudes toward privacy, trust levels, and comfort using commercial versus institutional generative AI tools.

The interviewees included two students—both non-STEM majors—and one staff member who had worked in both STEM and non-STEM colleges. Interviews lasted between 45 and 60 minutes and were conducted virtually. Participants were verbally informed before the interview that criticisms of the university or its tools would not impact their academic or professional standing, alleviating self-censorship concerns. We probed participants for specific incidents before asking for broader opinions in order to ground abstract concepts in concrete behaviors. We addressed sensitive topics such as policy noncompliance indirectly through hypothetical scenarios, which is the approach Wang et al. took [21].

The general outline for each interview began with questions about generative AI usage, such as participants’ preferred tools, first encounters, issues, attitudes, and their perceptions of how their peers use generative AI.

The next section prompted questions about interviewees’ views on privacy and security concerns regarding their preferred generative AI chatbots. Participants were encouraged to share what information they believed to be private, how they responded to privacy concerns, if any, and how they felt others responded to these issues.

The final section asked for opinions on privacy policies and guidance. Interviewees were asked if they had read privacy policies before, their experiences with generative AI policies and related risks, and what factors might influence a student’s decision to accept unclear data terms. Participants shared what tips and guidance have helped them and what UF could do to increase awareness of privacy risks.

These interviews were conducted remotely on Zoom and recorded and transcribed locally. Transcripts were cleaned to remove personally identifiable information. We verbally communicated these procedures to participants before the

interview. These interviews provided qualitative data that was analyzed using thematic coding.

4.2.3 Participants

The survey had a total of 57 respondents, with 89.5% (51) reported as students, 8.8% (5) reported as staff, and 1.8% (1) reported as “other.” There were no faculty respondents. Table 1 lists the demographic breakdown of participants. Of all respondents, 58% were female, with one participant identifying as non-binary or third gender. This result signals a need for a more inclusive participant pool to support security and privacy education that reaches diverse populations.

The age distribution clustered around 18 to 20 years, with a small tail extending to 27. As the sample reflected a predominantly young student body pursuing undergraduate degrees, there is a limit to the depth of experience of graduate-level students (6 total) who are likely handling research.

We organized reported academic majors into eight categories: Arts (Art, Music, Graphic Design), Computers and Math (Computer Science, Computer Engineering, Electrical Engineering, Electrical and Computer Engineering, Statistics, Biostatistics), Health (Applied Physiology and Kinesiology, Communications Sciences and Disorders, Health Education and Behavior, Nursing and Pre-nursing, Public Health), Humanities (English, History), Natural Sciences (Animal Sciences, Biology, Geology, Microbiology, Marine Science, Plant Molecular Cell Biology), Non-computers Engineering (Aerospace Engineering, Biomedical Engineering, Chemical Engineering, Environmental Engineering, Mechanical Engineering, Nuclear Engineering), Social Sciences (Criminology, PPEL (Philosophy, Politics, Economics, and Law), Political Science, Psychology), and Other (Agricultural Communication, Construction Management, Journalism).

4.2.4 Recruitment

To recruit participants for the survey, we distributed the link to individuals, group chats, computer science Discord servers, and student organizations. We targeted a variety of majors to understand the use of generative AI across campus—not limited to computer science or related fields. We encouraged respondents to share the survey with their peers. Faculty primarily received the survey via email. Interview requests were sent by email to a number of students, faculty, and staff, including the NaviGator team.

4.2.5 Ethics

We follow ethical research standards by ensuring voluntary participation and informed consent for all survey and interview participants. We required recruits to be at least 18 years old in order to participate. Recruits were informed of the study’s purpose, procedures, and their right to withdraw at any time. No personally identifiable data was collected and

Table 1: Demographics

Demographics	Student	Staff	Total
<i>Gender</i>			
Male	20 (40.0%)	0 (0%)	20 (36.4%)
Female	29 (58.0%)	5 (100.0%)	34 (61.8%)
Non-binary or third gender	1 (2.0%)	0 (0%)	1 (1.8%)
<i>Degree Program</i>			
Undergraduate	44 (88.0%)	0 (0%)	44 (88.0%)
Graduate	3 (6.0%)	0 (0%)	3 (6.0%)
Doctorate	3 (6.0%)	0 (0%)	3 (6.0%)

all data was anonymized and deleted after analysis. Data was stored on password-protected devices as stated by UF’s data protection guidelines.

5 Results

Listed below are the four research questions to interpret the survey and interview data:

RQ 1: How extensively do UF students, faculty, and staff use generative AI, and which tools do they use?

RQ 2: To what extent do people trust generative AI systems to handle their data responsibly and securely, and is trust different between commercial and institutional systems? What factors most prominently shape these perceptions? How do these perceptions of trust influence decisions in the use of generative AI?

RQ 3: What do people understand about how generative AI platforms handle data and privacy? How do people perceive the privacy protections of commercial platforms versus the models supported by NaviGator Chat, especially those labeled for sensitive data?

RQ 4: What are the privacy risks for populations within the university community in relation to generative AI systems? What should be considered when creating a generative AI system for higher education?

5.1 RQ1: Generative AI Usage Habits

To assess how extensively students, faculty, and staff at the University of Florida use generative AI, we examined whether respondents ever encountered the technology. As shown in Table 2, every survey participant had heard of a generative AI chatbot. Out of 57 survey responses, 72% (38) reported using some generative AI tool, but all respondents reported being at least “slightly familiar” with at least one of these systems. Note that although some individuals may have multiple roles, such as being a staff member and a student at the same time, each respondent in our survey only reported one role. All survey questions were optional, so the number of responses per question may not equal the sample size. The discussion

below focuses on a combined sample of students and staff since there were no faculty responses. Due to a low number of staff responses ($N = 5$), we cannot make strong conclusions about generative AI usage and privacy behaviors of UF staff.

Platform popularity (Table 2) reveals that the most popular generative AI chatbot platform among students, by a large margin, was OpenAI ChatGPT (87.5%), with Google Gemini taking second place (54.2%). The most popular chatbot platform among staff was tied between OpenAI ChatGPT and University of Florida NaviGator Chat at 60%, followed by Microsoft Copilot (40%). NaviGator (20.8%) and Copilot (14.8%) were not as popular among students, while Gemini was not as popular among staff (20%). Only three students reported using Claude. There may be a higher staff awareness of NaviGator Chat, by training or work requirement, that may explain the disparity of NaviGator Chat use between students and staff. Other generative AI platforms used in the past 6 months included Happy Scribe for transcription and Grammarly for editing.

The task distribution (Table 2) shows that the primary use of generative AI chatbots is seeking information, followed by writing or generating text. “Academic tasks or homework” and “academic research” were reported exclusively by students, which is consistent with existing research on university students. In contrast, staff used generative AI for “work tasks” and “productivity or organization,” which reflects the importance of office automation and project management aids. Other reported uses for generative AI included math, creating grocery lists, pantry-to-recipe ideas, programming, and assistance in editing papers.

From these results, we infer that a significant portion of the University of Florida community, including students, uses generative AI. Most students appear to have used AI chatbots recently since only 4.2% reported not using generative AI chatbots in the past 6 months.

Table 2: Usage of Generative AI by Role

Role	Student	Staff	Total
<i>Use of Generative AI</i>			
Used	35	3	38
Not used but heard of it	13	2	15
<i>Generative AI Platforms Used (Past 6 months)</i>			
OpenAI ChatGPT	42	3	45
Google Gemini	26	1	27
Anthropic Claude	3	0	3
Microsoft Copilot	7	2	9
UF NaviGator Chat	10	3	13
Other	0	2	2
No generative AI used	2	0	2
<i>Tasks</i>			
Seeking information	40	1	41
Writing or generating text	21	3	24
Academic tasks or homework	28	0	28
Academic research	16	0	16
Work tasks	14	2	16
Entertainment or self expression	7	1	8
Productivity or organization	17	2	19
Other	4	1	5

5.2 RQ2: Trust in Commercial Generative AI versus NaviGator Chat

Trust is the bridge between perceived risk and a user’s willingness to engage with a system. Previous research published by Haoyu Zhao et al. found that transparency and controllability significantly reduced perceived risks in generative AI systems and that women reported higher perceived risk [24]. When asked whether the respondent altered their behavior or avoided certain generative AI systems due to privacy or trust concerns, 43.3% of female respondents and 43.8% of male respondents answered yes, while 26.7% of female respondents and 31.3% of male respondents answered no. Male and female respondents altered their behavior at similar rates, but a higher number of male respondents did not alter their behavior.

Our survey also captured respondents’ level of trust in generative AI platforms with a five-point scale (1 = not at all, 5 = completely) that measured how much they trust each platform to handle data responsibly. Table 3 shows there appears to be a higher trust in ChatGPT among men, as 53.3% of women had little to no trust in ChatGPT compared to 41.1% of men. A gender difference was less apparent with Gemini, where 56.6% of women and 52.9% of men reported little to no trust in the system. For Claude and Microsoft Copilot, men were split between no trust and moderate trust. Claude performed the worst among women, with 78.1% reporting little to no trust. However, 52.1% of women and 32.5% of men had little

to no trust in NaviGator Chat.

Students generally rate their trust in chatbots higher than staff. The highest rating given to a platform by staff was “moderate” and only assigned to Microsoft Copilot and NaviGator Chat. Both groups rated NaviGator Chat the highest, but students chose ChatGPT as the runner-up (54.3% at “moderate” or above), while staff trusted Copilot over the rest (20% at “moderate”). However, although students had a higher trust in NaviGator Chat, it is important to note that not many have used this institutional platform. Only 25% of students reported using NaviGator Chat, and 31.3% reported they had never heard of it. In comparison, 60% of staff reported using Navigator Chat, but all staff respondents reported at least having heard of it. This finding was somewhat reflected in our interviews. Our student interviewees had never used NaviGator and reported no knowledge of the platform except that it exists. As a result, their perceptions of what NaviGator is, its purpose, and its data handling practices are based on speculation.

Meanwhile, responses to privacy risks varied. All five staff respondents reported that privacy or trust concerns influence how they use generative AI, while only 36.6% of students reported the same. The rest of the students were split between not altering behavior and avoiding use, or were not sure.

One of the interviewees stated that the intensity and scope of privacy are subjective. For example, certain types of information may or may not be considered private depending

Table 3: Trust in Platform Handling Data Responsibly by Gender

Trust Level	Male	Female	Non-binary	Total
<i>OpenAI ChatGPT</i>				
Not at all	4 (23.5%)	9 (30.0%)	1 (100%)	14 (29.2%)
A little	3 (17.6%)	7 (23.3%)	0 (0%)	10 (20.8%)
Moderately	8 (47.1%)	10 (33.3%)	0 (0%)	18 (37.5%)
A lot	2 (11.8%)	4 (13.3%)	0 (0%)	6 (12.5%)
Completely	0 (0%)	0 (0%)	0 (0%)	0 (0%)
<i>Google Gemini</i>				
Not at all	5 (29.4%)	10 (33.3%)	1 (100%)	16 (33.3%)
A little	4 (23.5%)	7 (23.3%)	0 (0%)	11 (22.9%)
Moderately	6 (35.3%)	11 (36.7%)	0 (0%)	17 (35.4%)
A lot	2 (11.8%)	2 (6.7%)	0 (0%)	4 (8.3%)
Completely	0 (0%)	0 (0%)	0 (0%)	0 (0%)
<i>Anthropic Claude</i>				
Not at all	6 (40%)	16 (57.1%)	1 (100%)	23 (52.3%)
A little	0 (0%)	6 (21.4%)	0 (0%)	6 (13.6%)
Moderately	8 (53.3%)	5 (17.9%)	0 (0%)	13 (29.5%)
A lot	1 (6.7%)	1 (3.6%)	0 (0%)	2 (4.5%)
Completely	0 (0%)	0 (0%)	0 (0%)	0 (0%)
<i>Microsoft Copilot</i>				
Not at all	7 (46.7%)	7 (24.1%)	1 (100%)	15 (33.3%)
A little	0 (0%)	10 (34.5%)	0 (0%)	10 (22.2%)
Moderately	7 (46.7%)	8 (27.6%)	0 (0%)	15 (33.3%)
A lot	1 (6.7%)	4 (13.8%)	0 (0%)	5 (11.1%)
Completely	0 (0%)	0 (0%)	0 (0%)	0 (0%)
<i>UF NaviGator Chat</i>				
Not at all	6 (20.0%)	3 (18.8%)	1 (100%)	10 (21.3%)
A little	2 (12.5%)	10 (33.3%)	0 (0%)	12 (25.5%)
Moderately	8 (50.0%)	10 (33.3%)	0 (0%)	18 (38.3%)
A lot	2 (12.5%)	3 (10.0%)	0 (0%)	5 (10.6%)
Completely	1 (6.3%)	1 (3.3%)	0 (0%)	2 (4.3%)

on personal privacy risks and disposition. We asked survey respondents how comfortable they were sharing certain types of information with NaviGator Chat compared to commercial alternatives. Staff were generally less comfortable sharing all types of data than students, but there was not much variability in comfort between the types of data. Students tend to be more uncomfortable sharing personal information, but many were “somewhat” or “very” comfortable sharing academic or professional work (41.6% and 38.9%), while only one staff member reported in the “somewhat comfortable” category, which was for professional work (Table 4).

Our results suggest that substantiated trust is a driver of generative AI use. Although students rated NaviGator Chat marginally higher than OpenAI ChatGPT in perceived responsible data handling, low familiarity and lack of understanding

keep usage low. For higher education institutions that wish to grow their generative AI tools, transparency (clear consent, public data-handling statements, accessible resources) and informed education on data protection practices are necessary. Enhancing these two factors may increase confidence in NaviGator as a more privacy-conscious alternative and, in turn, usage.

5.3 RQ3: Perception of Privacy and Data Handling in Commercial Generative AI versus NaviGator Chat

To address this research question, we asked survey respondents to rate generative AI platforms on transparency regarding privacy and data management practices. The purpose of

Table 4: Comfortability with sharing information on NaviGator Chat as opposed to commercial platforms

Comfort Level	Student	Staff	Total
<i>Personal information</i>			
Very uncomfortable	4 (10.8%)	2 (40.0%)	6 (14.3%)
Somewhat uncomfortable	15 (40.5%)	1 (20.0%)	16 (38.1%)
Neither comfortable nor uncomfortable	11 (29.7%)	2 (40.0%)	13 (31.0%)
Somewhat comfortable	5 (13.5%)	0 (0%)	5 (11.9%)
Very comfortable	2 (5.4%)	0 (0%)	2 (4.8%)
<i>Academic work</i>			
Very uncomfortable	3 (8.3%)	1 (20.0%)	4 (9.8%)
Somewhat uncomfortable	8 (22.2%)	2 (40.0%)	10 (24.4%)
Neither comfortable nor uncomfortable	10 (27.8%)	2. (40.0%)	12 (29.3%)
Somewhat comfortable	8 (22.2%)	0 (0%)	8 (19.5%)
Very comfortable	7.0 (19.4%)	0 (0%)	7 (17.1%)
<i>Professional work</i>			
Very uncomfortable	3 (8.3%)	2 (40.0%)	5 (12.2%)
Somewhat uncomfortable	7 (19.4%)	1 (20.0%)	8 (19.5%)
Neither comfortable nor uncomfortable	12 (33.3%)	1 (20.0%)	13 (31.7%)
Somewhat comfortable	9. (25.0%)	1 (20.0%)	10 (24.4%)
Very comfortable	5 (13.9%)	0 (0%)	5 (12.2%)

this was to understand how accessible privacy information is to users.

One of the interviewees speculated a correlation between academic major and trust because studying disciplines related to generative AI or privacy may impact awareness of privacy risks in generative AI. Computers and Math majors may be more familiar with how generative AI works whereas Social Sciences majors may be more familiar with technology ethics and privacy. Students in majors related to Computer Science and Math generally rated all chatbot platforms lower in trust and transparency compared to students in other fields of study, which might support the idea that greater awareness of generative AI privacy risks lowers trust in chatbot platforms for responsible data handling. Students in Computers and Math majors were the most critical of Copilot’s data handling transparency, with 71.4% rating Copilot “not at all” transparent. Non-computers Engineering and Natural Sciences majors generally rated transparency higher, with at least 50% rating moderate or higher. All students generally rated NaviGator Chat as more transparent than the other commercial options.

NaviGator Chat designates which models are appropriate for inputting private data under UF’s data classification labels. This feature allows users to use AI with UF sensitive or restricted data without violating UF’s data security policy. However, 41.7% of students report never paying attention to the data labels when using NaviGator Chat and 33.3% report

sometimes paying attention to them. There were only three staff respondents who used NaviGator Chat, but they reported paying attention to the data labels in varying frequencies. Since the sensitive and restricted labels only apply to data related to the university and its operations, many students may not notice or interact with data labels because they are not working with sensitive or restricted data.

Information about privacy management and data handling is usually found in a policy or terms of service that a platform asks users to read and agree to before granting access to the software product. However, few participants reported reading these documents for all generative AI chatbot platforms (9.3%), and 40% of staff and 60.5% of students reported not reading them. One interviewee also reported not reading the privacy policy or terms of service for the generative AI platforms they used, citing the text as too long and the language too technical for them to understand. Another interviewee believed that privacy policies and terms of service are “intentionally deceitful,” designed to make the user accept the platform’s terms without fully understanding them. This interviewee also remarked, “We’re so desensitized to giving up personal information to online platforms and not really caring about it. [...] [People] care about convenience; they want their homework done, [...] they’re thinking about their finals, their assignments. They’re not thinking about, ‘Oh what could this data be used for in 10, 20, 30 years?’” Based on this data, we

infer that respondents do not understand how generative AI platforms handle their data and privacy. Furthermore, interviewees stated that concern about data security often does not stop people from using generative AI platforms if others around them use them, if they believe the tool is reputable, or if the tool is convenient.

5.4 RQ4: Generative AI Risks and Considerations

The survey asked respondents to rate how strongly they felt about specific privacy concerns. These concerns are researching or training AI models on user data, monitoring or surveillance, sharing data with third parties, data compromise by security breach, and data leaks by AI, on a scale from “not at all” to “extremely.”

Staff reported a higher level of concern for all listed privacy risks than students (Table 5). No staff reported no concern for any privacy risk, and 60-80% reported being “extremely” concerned for each. The risk of monitoring or surveillance was the greatest concern among students, with 57.2% reporting being “very” or “extremely” concerned. 100% of staff were “very” or “extremely” concerned about model research or training on sensitive data, which was the lowest-ranking concern for students. In comparison, staff were least concerned about data compromised by a security breach. Since our staff sample is small, it’s inconclusive to say that staff rate every risk higher than students. However, working with others’ data and a higher awareness of how privacy and data security impact an organization may influence staff to rate privacy risks higher. Our staff interviewee pointed out that “[Students] only care if someone they don’t like has access to their data. People only care [about] what affects them directly.”

Our survey also queried how respondents attempt to mitigate the risks they are most concerned about. Table 6 summarizes mitigation measures by gender. The most common risk mitigation was limiting generative AI use cases. More men reported using different email addresses or personas than women (36.4% to 11.8%), but more women reported deciding not to use generative AI (41.2% to 9.1%). These mitigation strategies indicate that this population is responding to privacy threats but is limited by their understanding of privacy and what generative AI tools might be doing with their data.

These findings suggest that the messaging around NaviGator Chat is inconsistent in reaching the entire population NaviGator Chat serves. Staff, who appear to be more risk-aware, likely respond well to available documentation and training. However, students, whose primary concerns are surveillance and their data being shared, need clearer assurances about NaviGator Chat’s data collection and retention practices. Addressing knowledge gaps in the student body by educating them on online privacy, best practices, and mitigation techniques may foster trust in NaviGator Chat and increase usage rates.

6 Challenges

The primary challenge to the validity of our results is the small sample size. As the total number of respondents is only 57, with 89.5% of responses coming from students, we cannot reliably conclude how faculty and staff perceive generative AI. The same issue pertains to the student population because our sample size of 51 surveyors still risks generalization since the University of Florida’s total enrollment is 61,890 students as of Fall 2024 [15]. Due to the limited time frame of the course, our interview sample size was also low. We completed two student interviews and one staff interview to accompany our survey results. To mitigate these concerns as the project developed, we extended the survey and interview windows by a week. However, the question of the validity of our findings for the University of Florida population remains.

Another challenge questions the construct of our survey. The questions were inspired by previous work, observations from a staff retreat that held a course on NaviGator Chat and Microsoft Copilot, and our own inquiries about NaviGator Chat. We included an “other” field and a final feedback question to allow respondents to answer in ways we could not anticipate. Nevertheless, there remains an argument as to whether we asked the right questions due to our own biases, lack of perspective, and the absence of a pilot study.

Similarly, the accuracy of how well we collected and analyzed our survey results and interview transcripts poses a threat to the validity of our conclusions. The survey was rather long due to the scope of the research questions and the comparisons between commercial generative AI and NaviGator Chat, which risked respondent fatigue. To mitigate this issue, we implemented display logic to avoid asking questions that did not pertain to some respondents and randomized the order of options when possible.

7 Conclusion and Recommendations

We examined the perception of trust and transparency in commercial and institutional generative AI platforms at the University of Florida. Our results suggest that all students and staff are at least slightly familiar with generative AI platforms and that most have used one in some capacity in the past six months. With the growth of generative AI and the creation of NaviGator Chat, we anticipate that these tools and their popularity will only continue to expand to the point that generative AI literacy will become vital to society. As institutional generative AI platforms such as NaviGator Chat are still new and understudied, understanding how academic communities, especially students, trust and interact with these systems can help develop and improve generative AI literacy, privacy, and data management practices. This project can be used to (1) understand how students, faculty, and staff at the University of Florida perceive trust and transparency related to generative AI, (2) identify the risks of opaque design to create trustwor-

Table 5: Privacy Risks of Generative AI Usage

Concern Level	Student	Staff	Total
<i>Your data being used to research or train AI models</i>			
Not at all	4 (10.8%)	0 (0%)	4 (8.5%)
A little	17 (40.5%)	0 (0%)	17 (36.2%)
Moderately	8 (19.0%)	0 (0%)	8 (17.0%)
Very	7 (16.7%)	1 (20.0%)	8 (17.0%)
Extremely	6 (14.3%)	4 (80.0%)	10 (21.3%)
<i>Your activity being monitored or surveilled</i>			
Not at all	2 (4.8%)	0 (0%)	2 (4.3%)
A little	8 (19.0%)	0 (0%)	8 (17.0%)
Moderately	11 (26.2%)	1 (20.0%)	12 (25.5%)
Very	13 (26.2%)	1 (20.0%)	14 (29.8%)
Extremely	8 (19.0%)	3 (60.0%)	11 (23.4%)
<i>Your data being shared with third parties</i>			
Not at all	2 (4.8%)	0 (0%)	2 (4.3%)
A little	4 (9.5%)	0 (0%)	4 (8.5%)
Moderately	18 (42.9%)	2 (40.0%)	20 (42.6%)
Very	9 (21.4%)	0 (0%)	9 (19.1%)
Extremely	9 (21.4%)	3 (60.0%)	12 (25.5%)
<i>Your data being compromised in a security breach</i>			
Not at all	2 (4.8%)	0 (0%)	2 (4.3%)
A little	4 (9.5%)	1 (20%)	5 (10.6%)
Moderately	14 (33.3%)	1 (20.0%)	15 (31.9%)
Very	10 (23.8%)	0 (0%)	10 (21.3%)
Extremely	12 (28.6%)	3 (60.0%)	15 (31.9%)
<i>Your data being leaked by AI</i>			
Not at all	2 (4.8%)	0 (0%)	2 (4.3%)
A little	4 (9.5%)	1 (20%)	5 (10.6%)
Moderately	12 (28.6%)	0 (0%)	12 (25.5%)
Very	12 (28.6%)	0 (0%)	10 (25.5%)
Extremely	12 (28.6%)	4 (80.0%)	16 (34.0%)

thy and user-friendly systems for education, and (3) provide recommendations to improve generative AI literacy in higher education.

Despite widespread exposure to generative AI tools, our findings suggest that trust in these systems is based less on concrete knowledge of their function and more on perceptions of convenience, social norms, and assumptions about reputation. Respondents expressed general familiarity with generative AI systems, yet their understanding of data handling, privacy, and the underlying mechanisms remains limited. The gap between perceived trustworthiness and actual understanding stipulates that trust may be inflated relative to the transparency and data management practices of commercial systems.

Regarding NaviGator Chat, respondents expressed baseline

trust in the platform as an institutional tool, but their use of the platform was limited. Many were unaware of how NaviGator Chat differs from commercial platforms such as OpenAI ChatGPT and simply did not care to experiment because their preferred commercial platform(s) satisfied their needs. All three interviewees expressed concerns about monitoring and potential misuse when using NaviGator Chat. However, these concerns did not deter respondents from using generative AI altogether if platforms were endorsed by their peers, perceived as reputable, and produced satisfactory outputs.

This discrepancy between user trust, limited understanding, and opaque system designs highlights the need to develop transparent, secure, and ethical generative AI tools. NaviGator Chat is on the right path to promoting security and privacy within the university community, but the interface, poli-

Table 6: Privacy-risk mitigation actions by gender

Gender	Male	Female	Non-binary or Third Gender	Total
<i>Privacy Risk Mitigation Actions</i>				
Limit what tasks I use generative AI for	8 (72.7%)	12 (70.6%)	0 (0%)	20 (71.4%)
Omit, obfuscate or modify input	5 (45.5%)	7 (41.2%)	0 (0%)	12 (42.9%)
Provide false or misleading input	1 (9.1%)	2 (11.8%)	0 (0%)	3 (10.7%)
Apply privacy settings on generative AI tool (e.g., don't use my data for training the model)	3 (27.3%)	5 (29.4%)	0 (0%)	8 (28.6%)
Delete chat history	4 (36.4%)	5 (29.4%)	0 (0%)	9 (32.1%)
Choose certain generative AI tools to use	3 (27.3%)	4 (23.5%)	0 (0%)	7 (25.0%)
Use AI tools under different email addresses or personas	4 (36.4%)	2 (11.8%)	0 (0%)	6 (21.4%)
Use an external service or tool to mitigate privacy risk	2 (18.2%)	1 (5.9%)	0 (0%)	3 (10.7%)
Choose to not use generative AI at all	1 (9.1%)	7 (41.2%)	0 (0%)	8 (28.6%)
Other	1 (9.1%)	0 (0%)	0 (0%)	1 (3.6%)

cies, and accessibility need improvement. Therefore, based on these findings, we present three recommendations for improving generative AI trust and literacy in higher education:

- 1. Publish clear, accessible data-handling policies for institutional generative AI tools.**

We suggest that UF share how NaviGator Chat processes, stores, and protects user data in a straightforward manner. Similarly, data-label and retention policies should be easily accessible and explicitly stated in the policy documentation. For example, data-label descriptions could appear as a tooltip or a short pop-up when users hover over each label while selecting a model, rather than only displaying the label name. Transparent communication could reduce speculative concerns about monitoring and misuse.

- 2. Integrate privacy literacy and generative-AI best practices into the curriculum or orientation programming.**

Since UF staff already receive training on NaviGator Chat and Microsoft Copilot, and students are required to take violence-prevention and mental health courses, a comparable approach could be implemented for generative AI. Short courses, quizzes, or workshops on privacy, security, responsible use, and the limitations of generative AI would provide students with the foundational knowledge needed to use these tools appropriately. Such programs could increase the usage of NaviGator AI and foster trust by clarifying user expectations of the system's capabilities and constraints.

- 3. Develop tailored outreach to students, faculty, and staff individually.**

Since NaviGator Chat use cases vary by role, UF should customize communication and training for each group's needs and responsibilities. Highlighting NaviGator Chat's unique features and its security advantages over commercial tools may increase adoption while encouraging ongoing dialog about online safety, digital privacy, and ethical AI use.

These recommendations are a culmination of survey and interview responses intended to advance a campus environment where generative AI tools are used responsibly and ethically. As generative AI continues to evolve, institutions have the opportunity to guide their communities toward more informed, privacy-aware, and critical use of these technologies.

8 Acknowledgments

This project was conducted as a course assignment under the guidelines of *CIS 6930* and *CIS 4930 Security and Privacy for At-Risk Populations*. Any opinions, findings, conclusions, and recommendations shared in this paper are those of the authors and do not reflect the views of the University of Florida.

9 Appendix A: Survey Questions

AI Usage Questions

1. Do you use generative AI chat bot platforms (e.g. OpenAI ChatGPT, Google Gemini, Anthropic Claude, UF NaviGator Chat)?
 - ☐ No
 - ☐ Yes
2. How familiar are you with using generative AI chat bot platforms (e.g. OpenAI ChatGPT, Google Gemini, Anthropic Claude, UF NaviGator Chat)?
 - ☐ Not familiar at all
 - ☐ Slightly familiar
 - ☐ Moderately familiar
 - ☐ Very familiar
 - ☐ Extremely familiar
3. Which generative AI platforms have you used in the past 6 months? (Select all that apply.)
 - ☐ OpenAI ChatGPT
 - ☐ Google Gemini
 - ☐ Anthropic Claude
 - ☐ Microsoft Copilot
 - ☐ UF NaviGator Chat
 - ☐ Other
 - ☐ I have not used any generative AI chatbots in the past 6 months.
4. What do you typically use generative AI for? (Select all that apply.)
 - ☐ Seeking information
 - ☐ Writing or generating text
 - ☐ Academic tasks or homework
 - ☐ Academic research
 - ☐ Work tasks
 - ☐ Entertainment or self expression
 - ☐ Productivity or organization
 - ☐ Other
 - ☐ Not applicable
5. Have you used UF's NaviGator Chat?
 - ☐ Yes
 - ☐ No, but I've heard of it
 - ☐ No, but I haven't heard of it
6. Respondents who heard of NaviGator Chat: How did you hear about NaviGator Chat?
 - ☐ UF email
 - ☐ Work requirement
 - ☐ Training or workshop
 - ☐ Class
 - ☐ Word of mouth
 - ☐ Other
 - ☐ I don't remember
7. NaviGator Chat users: What do you use NaviGator Chat for? (select all that apply.)
 - ☐ Seeking information
 - ☐ Writing or generating text
 - ☐ Academic tasks or homework
 - ☐ Academic research
 - ☐ Work tasks
 - ☐ Entertainment or self expression
 - ☐ Productivity or organization
 - ☐ Other
8. NaviGator Chat users: Do you pay attention to the data labels (e.g. Open Data, Restricted Data, Sensitive Data) on NaviGator Chat?
 - ☐ Always
 - ☐ Most of the time
 - ☐ About half the time
 - ☐ Sometimes
 - ☐ Never
9. Respondents who did not respond *Never* to the previous question: Do the labels (e.g. Open Data, Restricted Data, Sensitive Data) influence your input on NaviGator Chat?
 - ☐ Strongly disagree
 - ☐ Somewhat disagree
 - ☐ Neither agree nor disagree
 - ☐ Somewhat agree
 - ☐ Strongly agree
10. Which of the following influence which AI platform(s), if any, you use (or don't use)? (Select all that apply.)
 - ☐ Recommendation from family, friends, peers
 - ☐ Requirement for a class or by UF
 - ☐ Public reputation
 - ☐ Privacy and transparency
 - ☐ Relevant and thorough output

- o Accurate and unbiased output
 - o Safe and non-harmful output
 - o Modes for content generation (e.g. image, audio, video)
 - o Accessibility features
 - o Customization features
 - o Integrations and connectors
 - o Other
11. Have you encountered difficulties while using generative AI tools?
- o Never
 - o Rarely
 - o Sometimes
 - o Frequently
 - o Always
 - o I haven't used a generative AI tool before

Privacy Questions

1. Have you or someone you know ever experienced unintended sharing of sensitive information or data leakage through AI-generated outputs?
 - o Yes
 - o No
2. Do you alter your behavior or avoid certain generative AI systems due to privacy or trust concerns?
 - o Yes
 - o No
 - o Not sure
3. Respondents who did not respond *No* to the previous question: Select the actions you do to mitigate privacy risks related to generative AI (select all that apply.)
 - o Limit what tasks I use generative AI for
 - o Omit, obfuscate or modify input
 - o Provide false or misleading input
 - o Apply privacy settings on generative AI tool (e.g., don't use my data for training the model)
 - o Delete chat history
 - o Choose certain generative AI tools to use
 - o Use AI tools under different email addresses or personas
 - o Use an external service or tool to mitigate privacy risk
 - o Choose to not use generative AI at all
 - o Other
4. NaviGator Chat users: Compared to commercial AI platforms (e.g. OpenAI ChatGPT, Google Gemini, Anthropic Claude), do you disclose more personal or sensitive information when using UF NaviGator Chat?
 - o Yes, I disclose more personal or sensitive information on NaviGator Chat
 - o No, I disclose more personal or sensitive information on commercial platforms
 - o No, my interactions are the same across platforms
 - o Not sure
 - o I only use UF Navigator Chat
5. Do you read the privacy policies or Terms of Service for any generative AI platform you have used or considered using?
 - o Yes, for all platforms
 - o Yes, for some platforms
 - o No
6. Select the statement that most closely aligns with your opinion. Compared to commercial generative AI platforms, NaviGator Chat has ...
 - o Significantly worse privacy protection
 - o Somewhat worse privacy protection
 - o Same level of privacy protection
 - o Somewhat better privacy protection
 - o Significantly better privacy protection

7. How much do you trust the following platforms to handle your data responsibly?

Trust Level	Not at all	A little	Moderately	A lot	Completely
OpenAI ChatGPT	o	o	o	o	o
Google Gemini	o	o	o	o	o
Anthropic Claude	o	o	o	o	o
Microsoft Copilot	o	o	o	o	o
UF NaviGator Chat	o	o	o	o	o

8. How transparent do you think each platform is about data handling and privacy practices?

Transparency Level	Not at all	A little	Moderately	A lot	Completely
OpenAI ChatGPT	o	o	o	o	o
Google Gemini	o	o	o	o	o
Anthropic Claude	o	o	o	o	o
Microsoft Copilot	o	o	o	o	o
UF NaviGator Chat	o	o	o	o	o

9. How concerned are you about the following privacy issues when using generative AI?

Concern Level	Not at all	A little	Moderately	A lot	Completely
Your data being used to research or train AI models	o	o	o	o	o
Your activity being monitored or surveilled	o	o	o	o	o
Your data being shared with third parties	o	o	o	o	o
Your data being compromised in a security breach	o	o	o	o	o
Your data being leaked by AI	o	o	o	o	o

10. How comfortable are you sharing the following information on UF NaviGator Chat compared to commercial platforms (e.g. OpenAI ChatGPT, Google Gemini, Anthropic Claude)?

Comfort Level	Very uncomfortable	Somewhat uncomfortable	Neither comfortable nor uncomfortable	Somewhat comfortable	Very comfortable
Personal information	o	o	o	o	o
Academic work	o	o	o	o	o
Professional work	o	o	o	o	o

Demographics Questions

1. What is your role at UF?
 - o Student
 - o Faculty
 - o Staff
 - o Other
2. Students: What is your major?
3. Students: Which degree program are you in?
 - o Bachelor's
 - o Master's
 - o Doctorate
4. Faculty: Which department are you a part of?
5. Staff: What is your job title?
6. How do you describe yourself?
 - o Male
 - o Female
 - o Non-binary / third gender
 - o Prefer to self-describe
 - o Prefer not to say
7. What is your age?
8. Choose one or more races that you consider yourself to be.
 - o White or Caucasian
 - o Black or African American
 - o American Indian/Native American or Alaska Native
 - o Asian
 - o Native Hawaiian or Other Pacific Islander
 - o Other
 - o Prefer not to say

References

- [1] AMOOZADEH, M., DANIELS, D., NAM, D., KUMAR, A., CHEN, S., HILTON, M., RAGAVAN, S. S., AND ALIPOUR, M. A. Trust in generative ai among students: An exploratory study. *SIGCSE 2024 1* (2024), 67–73.
- [2] BIN-NASHWAN, S. A., SADALLAH, M., AND BOUTERAA, M. Use of chatgpt in academia: Academic integrity hangs in the balance. *Technology in Society 75* (2023), 102370.
- [3] DORFMAN, E. What has uf been doing with ai so far? The Independent Florida Alligator, Sept. 2024.
- [4] DOTAN, R., PARKER, L. S., AND RADZILOWICZ, J. Responsible adoption of generative ai in higher education: Developing a “points to consider” approach based on faculty perspectives. *FACCT '24* (June 2024), 2033–2046.
- [5] GOLDA, A., MEKONEN, K., PANDEY, A., SINGH, A., HASSIJA, V., CHAMOLA, V., AND SIKDAR, B. Privacy and security concerns in generative ai: A comprehensive survey. *IEEE Access 12* (2024), 48126–48144.
- [6] HANCOCK, R. Uf unveils navigator ai, higher ed's first comprehensive, self-service ai platform. <https://news.ufl.edu/2024/08/navigator/>, 2024. Accessed 2025-12-08.
- [7] KORIR, M., SLADE, S., HOLMES, W., HÉLIOT, Y., AND RIENTIES, B. Investigating the dimensions of students' privacy concern in the collection, use and sharing of data for learning analytics. *Computers in Human Behavior Reports 9* (Mar. 2023), 100262.
- [8] LANKTON, N., MCKNIGHT, D., AND TRIPP, J. Technology, humanness, and trust: Rethinking trust in technology. *Journal of the Association for Information Systems 16* (10 2015), 880–918.
- [9] LAPPEMAN, J., MARLIE, S., JOHNSON, T., AND POGGENPOEL, S. Trust and digital privacy: willingness to disclose personal information to banking chatbot services. *Journal of Financial Services Marketing 28*, 2 (2023), 337–357.
- [10] MAHMOUD, A. B., KUMAR, V., AND SPYROPOULOU, S. Identifying the public's beliefs about generative artificial intelligence: A big data approach. *IEEE Transactions on Engineering Management 72* (2025), 827–841.
- [11] MAYER, R. C., DAVIS, J. H., AND SCHOORMAN, F. D. An integrative model of organizational trust. *The Academy of Management Review 20*, 3 (1995), 709–734.
- [12] OF FLORIDA, U. Bring ai to your data securely. <https://ai.ufl.edu/teaching-with-ai/for-uf-faculty/navigator/>, 2024. Accessed 2025-12-08.
- [13] OF FLORIDA, U. Data classification policy. <https://policy.ufl.edu/policy/data-classification-policy/>, 2024. Accessed 2025-12-08.
- [14] OF FLORIDA, U. Online/internet privacy policy. <https://privacy.ufl.edu/policies/onlineinternet-privacy-policy-/>, 2024. Accessed 2025-12-08.
- [15] OF FLORIDA, U. Uf facts. <https://news.ufl.edu/for-media/uf-facts/>, 2024. Accessed 2025-12-08.
- [16] OF FLORIDA, U. Ai across the curriculum. <https://ai.ufl.edu/teaching-with-ai/ai-across-the-curriculum/>, 2025. Accessed 2025-12-08.
- [17] RANE, N. Enhancing the quality of teaching and learning through chatgpt and similar large language models: Challenges, future prospects, and ethical considerations in education. *TESOL and Technology Studies 5*, 1 (Feb. 2024), 1–6.
- [18] RAO, H. Ethical and legal considerations behind the prevalence of chatgpt: Risks and regulations. *Frontiers in Computing and Intelligent Systems 4*, 1 (June 2023), 23–29.

- [19] SAÚDE, S., BARROS, J. P., AND ALMEIDA, I. Impacts of generative artificial intelligence in higher education: Research trends and students' perceptions. *Social Sciences* 13, 8 (2024).
- [20] SOUSA, S., CRAVINO, J., AND MARTINS, P. Challenges and trends in user trust discourse in ai popularity. *Multimodal Technologies and Interaction* 7, 2 (2023).
- [21] WANG, N., LI, Y., AND CONG, F. University students' privacy concerns toward generative artificial intelligence. *Journal of Academic Ethics* 23, 12 (June 2025), 2401–2422.
- [22] WU, X., DUAN, R., AND NI, J. Unveiling security, privacy, and ethical concerns of chatgpt. *Journal of Information and Intelligence* 2, 2 (Mar. 2024), 102–115.
- [23] YU, Y., LIU, Y., ZHANG, J., HUANG, Y., AND WANG, Y. Understanding generative ai risks for youth: A taxonomy based on empirical data. *arXiv preprint arXiv:2502.16383* (2025).
- [24] ZHAO, H., HAN, Z., YIN, S., YANG, N., AND HANSEN, P. From interface to inference: mapping the impact of generative artificial intelligence affordances on user risk perception. *Telematics and Informatics* 101 (2025), 102299.
- [25] ZHAO, H., YUAN, B., AND SONG, Y. Employees' perception of generative artificial intelligence and the dark side of work outcomes. *Journal of Hospitality and Tourism Management* 61 (Dec. 2024), 191–199.